

LLM Transformer Decoder RTL Accelerator — Technical Report

1. Introduction

This report documents the design, implementation, and verification of a synthesizable RTL accelerator for LLM (Large Language Model) transformer decoder inference. The accelerator implements a single transformer decoder block — the fundamental repeating unit of GPT-style language models — in SystemVerilog, targeting FPGA and ASIC deployment.

The design provides two architectural variants: a **register-bridge** variant that copies BRAM weights into register arrays for single-cycle combinational access (maximum throughput), and a **streaming** variant that reads weights directly from BRAM one element per cycle (minimum area). Both share the same external interface, BRAM storage, and verification infrastructure.

Key design features include Q8.8 fixed-point arithmetic, systolic array compute, BRAM-backed weight and KV-cache storage, division-free softmax normalisation via reciprocal LUT with Newton-Raphson refinement, and packed array ports for full simulation visibility.

2. Background

2.1 Transformer Decoder Architecture

The transformer decoder processes tokens autoregressively: given a sequence of prior tokens, it predicts the next token. Each decoder layer applies two sub-layers with residual connections. First, Multi-Head Self-Attention computes attention scores between the current query and all prior key/value pairs, enabling the model to attend to relevant context. Second, a Feed-Forward Network (FFN) applies a two-layer MLP with non-linear activation independently to each position. Both sub-layers are preceded by Layer Normalization (pre-norm architecture) and followed by residual additions.

2.2 Inference vs Training

This accelerator targets inference only, which has distinct characteristics compared to training. Batch sizes are typically 1 for interactive generation. Only the forward pass is needed with no backpropagation. KV-cache eliminates redundant attention computation. Weights are fixed, enabling weight-stationary dataflows and BRAM preloading.

2.3 Fixed-Point Quantization

Modern LLM inference widely uses quantization to reduce memory bandwidth and compute cost. Our Q8.8 fixed-point format (8 integer bits, 8 fractional bits) provides a range of -128.0 to +127.996 with 1/256 resolution, which is sufficient for demonstrating the architectural concepts while keeping multiplier width at 16 bits.

3. Architecture

3.1 Top-Level Block Diagram

The design has two levels of hierarchy. The synthesis top-level wraps the compute core with BRAM-backed storage for all 49,344 weight parameters and 16,384 KV-cache entries. Two top-level variants exist: `transformer_decoder_top` (register-bridge) and `transformer_decoder_top_stream` (direct BRAM streaming).

The compute core implements the pre-norm architecture used by GPT-2, LLaMA, and most modern LLMs: LayerNorm 1, Multi-Head Attention with KV-cache, residual addition, LayerNorm 2, Feed-Forward Network with ReLU, and a second residual addition.

3.2 Module Hierarchy

The full project comprises 17 RTL source files in two parallel hierarchies:

Register-bridge variant (original — maximum throughput):

```
transformer_decoder_top          (BRAM → register arrays → decoder)
  └── bram_sp ×12              (Weight/bias/parameter storage)
  └── kv_cache_bram ×2         (K and V caches)
    └── bram_dp                (True dual-port BRAM)
  └── transformer_decoder       (Decoder compute core)
    ├── layer_norm ×2          (Pre-attention & pre-FFN normalisation)
    └── multi_head_attention    (Causal multi-head self-attention)
```

└ softmax_unit	(Reciprocal-LUT softmax, time-multiplexed)
└ feed_forward	(Two-layer FFN with ReLU)
└ systolic_array	(Matrix multiply engine)
└ processing_element	(Single MAC unit)
└ transformer_pkg	(Parameters, types, FP utilities)

Streaming variant (new — minimum area):

transformer_decoder_top_stream	(BRAMs → address/data → decoder)
└ bram_sp ×12	(Same BRAMs, muxed read ports)
└ kv_cache_bram ×2	(Element-level BRAM reads)
└ bram_dp	
└ transformer_decoder_stream	(Streaming decoder core)
└ layer_norm ×2	(Unchanged – small packed ports)
└ multi_head_attention_stream	(BRAM addr/data weight reads)
└ softmax_unit	(Unchanged)
└ feed_forward_stream	(BRAM addr/data weight reads)
└ systolic_array	(Shared infrastructure)
└ processing_element	
└ transformer_pkg	

Shared modules (transformer_pkg, bram_sp, bram_dp, kv_cache_bram, softmax_unit, layer_norm, systolic_array, processing_element) are identical between variants. The streaming variant adds 4 new modules: multi_head_attention_stream, feed_forward_stream, transformer_decoder_stream, and transformer_decoder_top_stream.

3.3 Module Descriptions

3.3.1 Processing Element (PE)

The PE is the atomic compute unit — a single multiply-accumulate (MAC) cell. It receives an activation from the left and a weight from the top, computes $\text{acc} += \text{a} \times \text{w}$ using full 32-bit precision, and forwards both operands to neighboring PEs. The 32-bit accumulator prevents overflow during long dot-product chains. Key properties include single-cycle MAC latency, registered inputs and outputs for timing closure, synchronous clear for accumulator reset between tiles, and parameterized data widths via transformer_pkg.

3.3.2 Systolic Array

A 4×4 grid of PEs forms the systolic matrix multiplication engine. Data flows through the array in a wave pattern: activations propagate left-to-right and weights top-to-bottom. After ROWS + COLS – 1 cycles of

streaming, the 16 accumulators hold the complete output tile. This architecture achieves 16 MACs per cycle with only 8 input operands, giving a 2 \times data reuse factor.

3.3.3 Softmax Unit (Division-Free)

Softmax normalisation computes $\text{probs}[i] = \exp(\text{scores}[i] - \max) / \sum \exp$. The implementation proceeds through six FSM stages: find maximum score for numerical stability, subtract maximum and apply PWL exponential approximation, accumulate the sum of exponentials, compute the reciprocal of the sum via LUT + Newton-Raphson, and normalise by multiplying each exponential by the reciprocal.

The PWL exponential uses four linear segments covering [-8, 0] with non-negative clamping at segment boundaries. The normalisation step uses a 32-entry reciprocal LUT (indexed by the top 5 mantissa bits of the CLZ-normalised sum) providing an initial estimate in Q2.14. One Newton-Raphson iteration doubles precision to approximately 12 bits. The entire reciprocal computation is combinational (single-cycle) using a 512-bit ROM, two 16 \times 16 multipliers, and a barrel shifter. Maximum error versus exact division is ± 1 LSB.

3.3.4 Layer Normalization

LayerNorm computes the mean and variance of the input vector, then normalizes each element through three FSM stages: mean computation via sequential accumulation, variance computation using centered differences, and element-wise normalization with learnable gamma and beta parameters. The reciprocal square root uses a 4-entry LUT.

3.3.5 Multi-Head Attention

The attention module implements multi-head scaled dot-product attention with integrated softmax and KV-cache support. A single softmax_unit (VEC_LEN = 128) is time-multiplexed across N_HEADS = 4 heads. The register-bridge variant receives full weight matrices on array ports. The streaming variant replaces these with BRAM address/data interfaces, reading Wq/Wk/Wv in lockstep with three parallel accumulators.

3.3.6 Feed-Forward Network

The FFN implements $\text{FFN}(x) = \text{ReLU}(x \times W1 + b1) \times W2 + b2$. The register-bridge variant computes full dot products per cycle. The streaming variant reads $W1$, $b1$, and $W2$ from BRAM one element per cycle. Both share identical ReLU logic.

3.3.7 BRAM Modules

Three BRAM types: `bram_sp` (single-port, read-first, optional hex init), `bram_dp` (true dual-port), and `kv_cache_bram` (vector-write FSM + element-level read).

4. Streaming Architecture

4.1 Motivation

The register-bridge architecture copies all 49,344 BRAM words into register arrays, consuming approximately 789,504 flip-flops. The streaming architecture eliminates these by having compute modules generate BRAM read addresses directly from their FSM state.

4.2 Weight Read Protocol

Each streaming module exposes address/data/enable ports for its weight BRAMs. The top-level muxes each BRAM's read port between the weight-loading bus and the compute address. BRAM latency is one cycle: address issued in cycle N, data arrives in cycle N+1.

For QKV projection, Wq , Wk , and Wv share the same read address. Three accumulators run simultaneously, producing Q , K , V for one dimension in D_{MODEL} cycles.

4.3 Performance Comparison

Metric	Register-Bridge	Streaming	Change
Total pipeline (seq_pos=0)	~3,200 cycles	~45,000 cycles	~14× slower
Weight registers	789K FFs	5K FFs	99.4% reduction
Weight multipliers	~192 combinational	6 sequential	97% reduction

Metric	Register-Bridge	Streaming	Change
DSP48 usage	~150	~10	93% reduction

5. Verification

5.1 Summary

Total: **83/83 tests pass** (54 behavioural + 29 RTL simulation across 6 testbenches).

5.2 Behavioural Model (54 tests)

Fixed-Point Utilities (16/16), Processing Element (9/9), Systolic Array (8/8), Softmax Unit (8/8), Layer Normalization (5/5), Feed-Forward Network (4/4), Decoder Integration (12/12).

5.3 RTL Simulation (29 tests)

Processing Element (6/6), Systolic Array (4/4), Softmax (4/4), BRAM (8/8), Decoder — Register-Bridge (4/4), Decoder — Streaming (3/3).

The streaming decoder testbench loads weights via the `wl_en` bus, then verifies single-token and sequential two-token inference. Output `out_emb[0] = 0x00A8` (0.656) matches the register-bridge variant's first dimension.

5.4 RTL Lint

17 RTL files: 0 errors, 5 warnings (BRAM blocks intentionally omit reset).

6. Implementation Metrics (Estimated)

Block	DSP48	FFs	LUTs	BRAM18K
Processing Element	1	~30	~20	—
4x4 Systolic Array	16	~500	~350	—
Softmax Unit	2	~2K	~3K	—
Layer Normalisation	2–4	~300	~500	—
Multi-Head Attention	~64	~5K	~8K	—

Block	DSP48	FFs	LUTs	BRAM18K
Feed-Forward Network	~64	~3K	~5K	—
Weight BRAMs (12)	—	—	—	~110
KV-Cache BRAMs (2)	—	—	—	~16
Full Decoder (register-bridge)	~150	~11K	~17K	~126
Full Decoder (streaming)	~10	~5K	~12K	~126

Clock frequency estimate: 100–200 MHz.

7. Limitations and Future Work

7.1 Resolved

BRAM-backed weights, division-free softmax, packed array ports, streaming weight architecture (99.4% FF reduction).

7.2 Remaining

LayerNorm still uses $/$ operator and coarse rsqrt LUT. Only ReLU activation. Systolic array not connected to datapath. No batching.

7.3 Planned

LayerNorm right-shift + reciprocal-LUT, systolic-tiled projections, parallel softmax, GELU/SiLU PWL, multi-layer stacking, multi-device distribution, AXI-Lite interface.

8. Conclusion

This project demonstrates a complete, synthesizable transformer decoder block in SystemVerilog with two architectural variants: a high-throughput register-bridge design and a minimum-area streaming design. The 17-module hierarchy mirrors the transformer's conceptual structure while addressing practical synthesis concerns. All 83 verification tests pass across both variants. The streaming architecture achieves 99.4% register reduction while maintaining functional equivalence, making the design feasible on the smallest FPGA targets.